

Final Project: Reinforcement Learning for Autonomous Driving in Highway-Env

Course: CSC 580 - Artificial Intelligence II

Group-10

Members:

- Sanjay Uppala (2164383)
- Yogitha Pollisetty (2130220)

DePaul University

Installing required poackages

```
In [1]: !pip install highway-env
!pip install stable-baselines3
!pip install gymnasium
!pip install opencv-python
!pip install tensorboard
```

Requirement already satisfied: highway-env in /opt/anaconda3/lib/python3.12/site-packages (1.10.1)

Requirement already satisfied: gymnasium>=1.0.0a2 in /opt/anaconda3/lib/python3.12/site-packages (from highway-env) (1.0.0)

Requirement already satisfied: farama-notifications>=0.0.1 in /opt/anaconda3/lib/python3.12/site-packages (from highway-env) (0.0.4)

Requirement already satisfied: numpy>=1.21.0 in /opt/anaconda3/lib/python3.12/site-packages (from highway-env) (1.26.4)

Requirement already satisfied: pygame>=2.0.2 in /opt/anaconda3/lib/python3.12/site-packages (from highway-env) (2.6.1)

Requirement already satisfied: matplotlib in /opt/anaconda3/lib/python3.12/site-packages (from highway-env) (3.9.2)

Requirement already satisfied: pandas in /opt/anaconda3/lib/python3.12/site-packages (from highway-env) (2.2.2)

Requirement already satisfied: scipy in /opt/anaconda3/lib/python3.12/site-packages (from highway-env) (1.13.1)

Requirement already satisfied: cloudpickle>=1.2.0 in /opt/anaconda3/lib/python3.12/site-packages (from gymnasium>=1.0.0a2->highway-env) (3.0.0)

Requirement already satisfied: typing-extensions>=4.3.0 in /opt/anaconda3/lib/python3.12/site-packages (from gymnasium>=1.0.0a2->highway-env) (4.11.0)

Requirement already satisfied: contourpy>=1.0.1 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->highway-env) (1.2.0)

Requirement already satisfied: cycler>=0.10 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->highway-env) (0.11.0)

Requirement already satisfied: fonttools>=4.22.0 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->highway-env) (4.51.0)

Requirement already satisfied: kiwisolver>=1.3.1 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->highway-env) (1.4.4)

Requirement already satisfied: packaging>=20.0 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->highway-env) (24.1)

Requirement already satisfied: pillow>=8 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->highway-env) (10.4.0)

Requirement already satisfied: pyparsing>=2.3.1 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->highway-env) (3.1.2)

Requirement already satisfied: python-dateutil>=2.7 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->highway-env) (2.9.0.post0)

Requirement already satisfied: pytz>=2020.1 in /opt/anaconda3/lib/python3.12/site-packages (from pandas->highway-env) (2024.1)

Requirement already satisfied: tzdata>=2022.7 in /opt/anaconda3/lib/python3.12/site-packages (from pandas->highway-env) (2023.3)

Requirement already satisfied: six>=1.5 in /opt/anaconda3/lib/python3.12/site-packages (from python-dateutil>=2.7->matplotlib->highway-env) (1.16.0)

Requirement already satisfied: stable-baselines3 in /opt/anaconda3/lib/python3.12/site-packages (2.5.0)

Requirement already satisfied: gymnasium<1.1.0,>=0.29.1 in /opt/anaconda3/lib/python3.12/site-packages (from stable-baselines3) (1.0.0)

Requirement already satisfied: numpy<3.0,>=1.20 in /opt/anaconda3/lib/python3.12/site-packages (from stable-baselines3) (1.26.4)

Requirement already satisfied: torch<3.0,>=2.3 in /opt/anaconda3/lib/python3.12/site-packages (from stable-baselines3) (2.6.0)

Requirement already satisfied: cloudpickle in /opt/anaconda3/lib/python3.12/site-packages (from stable-baselines3) (3.0.0)

Requirement already satisfied: pandas in /opt/anaconda3/lib/python3.12/site-packages (from stable-baselines3) (2.2.2)

Requirement already satisfied: matplotlib in /opt/anaconda3/lib/python3.12/site-packages (from stable-baselines3) (3.9.2)

Requirement already satisfied: typing-extensions>=4.3.0 in /opt/anaconda3/lib/python3.12/site-packages (from gymnasium<1.1.0,>=0.29.1->stable-baselines3) (4.11.0)

Requirement already satisfied: farama-notifications>=0.0.1 in /opt/anaconda3/lib/python3.12/site-packages (from gymnasium<1.1.0,>=0.29.1->stable-baselines3) (0.0.4)

Requirement already satisfied: filelock in /opt/anaconda3/lib/python3.12/site-packages (from torch<3.0,>=2.3->stable-baselines3) (3.13.1)

Requirement already satisfied: networkx in /opt/anaconda3/lib/python3.12/site-packages (from torch<3.0,>=2.3->stable-baselines3) (3.3)

Requirement already satisfied: jinja2 in /opt/anaconda3/lib/python3.12/site-packages (from torch<3.0,>=2.3->stable-baselines3) (3.1.4)

Requirement already satisfied: fsspec in /opt/anaconda3/lib/python3.12/site-packages (from torch<3.0,>=2.3->stable-baselines3) (2024.6.1)

Requirement already satisfied: setuptools in /opt/anaconda3/lib/python3.12/site-packages (from torch<3.0,>=2.3->stable-baselines3) (75.1.0)

Requirement already satisfied: sympy==1.13.1 in /opt/anaconda3/lib/python3.12/site-packages (from torch<3.0,>=2.3->stable-baselines3) (1.13.1)

Requirement already satisfied: mpmath<1.4,>=1.1.0 in /opt/anaconda3/lib/python3.12/site-packages (from sympy==1.13.1->torch<3.0,>=2.3->stable-baselines3) (1.3.0)

Requirement already satisfied: contourpy>=1.0.1 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->stable-baselines3) (1.2.0)

Requirement already satisfied: cycler>=0.10 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->stable-baselines3) (0.11.0)

Requirement already satisfied: fonttools>=4.22.0 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->stable-baselines3) (4.51.0)

Requirement already satisfied: kiwisolver>=1.3.1 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->stable-baselines3) (1.4.4)

Requirement already satisfied: packaging>=20.0 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->stable-baselines3) (24.1)

Requirement already satisfied: pillow>=8 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->stable-baselines3) (10.4.0)

Requirement already satisfied: pyparsing>=2.3.1 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->stable-baselines3) (3.1.2)

Requirement already satisfied: python-dateutil>=2.7 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib->stable-baselines3) (2.9.0.post0)

Requirement already satisfied: pytz>=2020.1 in /opt/anaconda3/lib/python3.12/site-packages (from pandas->stable-baselines3) (2024.1)

Requirement already satisfied: tzdata>=2022.7 in /opt/anaconda3/lib/python3.12/site-packages (from pandas->stable-baselines3) (2023.3)

Requirement already satisfied: six>=1.5 in /opt/anaconda3/lib/python3.12/site-packages (from python-dateutil>=2.7->matplotlib->stable-baselines3) (1.16.0)

Requirement already satisfied: MarkupSafe>=2.0 in /opt/anaconda3/lib/python3.12/site-packages (from jinja2->torch<3.0,>=2.3->stable-baselines3) (2.1.3)

Requirement already satisfied: gymnasium in /opt/anaconda3/lib/python3.12/site-packages (1.0.0)

Requirement already satisfied: numpy>=1.21.0 in /opt/anaconda3/lib/python3.12/site-packages (from gymnasium) (1.26.4)

Requirement already satisfied: cloudpickle>=1.2.0 in /opt/anaconda3/lib/python3.12/site-packages (from gymnasium) (3.0.0)

Requirement already satisfied: typing-extensions>=4.3.0 in /opt/anaconda3/lib/python3.12/site-packages (from gymnasium) (4.11.0)

Requirement already satisfied: farama-notifications>=0.0.1 in /opt/anaconda3/lib/python3.12/site-packages (from gymnasium) (0.0.4)

Requirement already satisfied: opencv-python in /opt/anaconda3/lib/python3.12/site-packages (4.11.0.86)

Requirement already satisfied: numpy>=1.21.2 in /opt/anaconda3/lib/python3.12/site-packages (from opencv-python) (1.26.4)

Requirement already satisfied: tensorboard in /opt/anaconda3/lib/python3.12/site-packages (2.19.0)

Requirement already satisfied: absl-py>=0.4 in /opt/anaconda3/lib/python3.12/site-packages (from tensorboard) (2.1.0)

Requirement already satisfied: grpcio>=1.48.2 in /opt/anaconda3/lib/python3.12/site-packages (from tensorboard) (1.71.0)

Requirement already satisfied: markdown>=2.6.8 in /opt/anaconda3/lib/python3.12/site-packages (from tensorboard) (3.4.1)

Requirement already satisfied: numpy>=1.12.0 in /opt/anaconda3/lib/python3.12/site-packages (from tensorboard) (1.26.4)

Requirement already satisfied: packaging in /opt/anaconda3/lib/python3.12/site-packages (from tensorboard) (24.1)

Requirement already satisfied: protobuf!=4.24.0,>=3.19.6 in /opt/anaconda3/lib/python3.12/site-packages (from tensorboard) (4.25.3)
Requirement already satisfied: setuptools>=41.0.0 in /opt/anaconda3/lib/python3.12/site-packages (from tensorboard) (75.1.0)
Requirement already satisfied: six>1.9 in /opt/anaconda3/lib/python3.12/site-packages (from tensorboard) (1.16.0)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /opt/anaconda3/lib/python3.12/site-packages (from tensorboard) (0.7.2)
Requirement already satisfied: werkzeug>=1.0.1 in /opt/anaconda3/lib/python3.12/site-packages (from tensorboard) (3.0.3)
Requirement already satisfied: MarkupSafe>=2.1.1 in /opt/anaconda3/lib/python3.12/site-packages (from werkzeug>=1.0.1->tensorboard) (2.1.3)

```
In [5]: import os
import time
import glob
import io
import base64
import numpy as np
import gymnasium as gym
import highway_env
import cv2
from stable_baselines3 import DQN, A2C
from stable_baselines3.common.monitor import Monitor
from stable_baselines3.common.evaluation import evaluate_policy
from IPython import get_ipython
from IPython.display import HTML, display
from IPython import get_ipython
try:
    IN_JUPYTER = True
except ImportError:
    IN_JUPYTER = False
```

2) Directory Setup

```
In [3]: PROJECT_DIR = os.path.join(os.getcwd(), "Project_A2C")
LOG_DIR = os.path.join(PROJECT_DIR, "logs")
MODEL_DIR = os.path.join(PROJECT_DIR, "models")
VIDEO_DIR = os.path.join(PROJECT_DIR, "videos")

for folder in [PROJECT_DIR, LOG_DIR, MODEL_DIR, VIDEO_DIR]:
    os.makedirs(folder, exist_ok=True)
```

3) Utility: show_video in Jupyter

```
In [4]: def show_video(video_path_pattern):
    ...
    ... Display the latest video in the Notebook
    ...
    mp4_list = glob.glob(video_path_pattern)
    if not mp4_list:
        print(f"[WARN] No video found for: {video_path_pattern}")
        return
    video_path = mp4_list[0]
    with open(video_path, "rb") as f:
        video_data = f.read()
```

```

encoded_video = base64.b64encode(video_data).decode("ascii")
if IN_JUPYTER:
    display(HTML(data=f'''
        <video width="640" height="480" controls>
            <source src="data:video/mp4;base64,{encoded_video}" type="video/mp4">
        </video>
    '''))
else:
    print(f"[INFO] Video saved: {video_path}")

```

4) Environment Creation

```

In [5]: def make_env(env_name="highway-fast-v0"):
    """
    Created a Configure for the Enviornment using Gym for traning
    """
    env = gym.make(env_name, render_mode="rgb_array")
    if hasattr(env.unwrapped, "configure"):
        env.unwrapped.configure({
            "lanes_count": 4,
            "vehicles_count": 20,
            "duration": 40,
            "collision_reward": -2,
            "high_speed_reward": 1,
            "reward_speed_range": [20, 30],
            "action": {"type": "DiscreteMetaAction"}
        })
    return Monitor(env)

```

5) Video Recording with OpenCV

```

In [6]: def record_with_opencv(frames, out_file="demo.mp4", fps=15):
    """
    OpennCv package has been used to record the video in 3 format-MP4,AVC1,XVID
    """
    if not frames:
        print("[WARN] No frames to record.")
        return
    height, width, _ = frames[0].shape
    fourcc = cv2.VideoWriter_fourcc(*'avc1')
    out = cv2.VideoWriter(out_file, fourcc, fps, (width, height))
    for frame in frames:
        bgr_frame = cv2.cvtColor(frame, cv2.COLOR_RGB2BGR)
        out.write(bgr_frame)
    out.release()
    print(f"[INFO] Created video: {out_file}")

```

6) Training A2C Model

```

In [7]: if __name__ == "__main__":
    ENVIRONMENTS = ["highway-fast-v0", "intersection-v0", "roundabout-v0"]
    LEARNING_RATES = [0.0005, 0.001]
    GAMMAS = [0.95, 0.8]

    TOTAL_TIMESTEPS = 100000

```

```

# dictionary to store final results
results = {}

for env_name in ENVIRONMENTS:
    for lr in LEARNING_RATES:
        for gm in GAMMAS:
            # Unique run_id
            run_id = f"{env_name}_A2C_lr{lr}_gm{gm}"
            print(f"TRAINING RUN: {run_id}")

            # Create environment
            vec_env = make_env(env_name)

            tb_log_dir = os.path.join(LOG_DIR, run_id)

            # Initialize the A2C model
            model = A2C(
                "MlpPolicy", # Multi-layer Perceptron policy
                vec_env,      # Gym environment
                learning_rate=lr,
                gamma=gm,     # Discount factor
                n_steps=5,    # Number of steps before updating gradients
                vf_coef=0.5,  # Value function coefficient
                ent_coef=0.01, # Entropy regularization (improves exploration)
                verbose=0,
                tensorboard_log=tb_log_dir
            )

            # Train
            model.learn(total_timesteps=TOTAL_TIMESTEPS)

            # Evaluate
            mean_reward, std_reward = evaluate_policy(model, vec_env, n_eval_episode)

            if mean_reward > 0:
                success_rate = (mean_reward - std_reward) / mean_reward
            else:
                success_rate = 0

            print(f" Completed run for env={env_name}, gamma={gm}, lr={lr}")
            print(f"   Final average reward: {mean_reward:.2f}")
            print(f"   Success rate: {success_rate:.2f} \n")

            # Save model
            model_path = os.path.join(MODEL_DIR, f"{run_id}_model")
            model.save(model_path)
            print(f"[INFO] Model saved at: {model_path}\n")

            # Record short video
            demo_env = make_env(env_name)
            if hasattr(demo_env.unwrapped, "config"):
                demo_env.unwrapped.config["simulation_frequency"] = 15

            frames = []
            obs, info = demo_env.reset()
            done = truncated = False
            for step in range(300):
                action, _ = model.predict(obs, deterministic=True)
                obs, reward, done, truncated, info = demo_env.step(action)
                frames.append(demo_env.render())
                if done or truncated:

```

```

        obs, info = demo_env.reset()
        demo_env.close()

    video_file = os.path.join(VIDEO_DIR, f"{run_id}_{int(time.time())}.mp4")
    record_with_opencv(frames, out_file=video_file, fps=15)
    show_video(video_file)

    results[run_id] = {
        "env": env_name,
        "lr": lr,
        "gamma": gm,
        "mean_reward": mean_reward,
        "std_reward": std_reward,
        "success_rate": success_rate
    }

    # Display TENSORBOARD plot for this run using logs file
    if IN_JUPYTER:
        get_ipython().run_line_magic('load_ext', 'tensorboard')
        get_ipython().run_line_magic('tensorboard', '--logdir ./Project_A2C/')

    print("\nFinal Result of All model trained Above\n")
    for i in results:
        print(i)

```

TRAINING RUN: highway-fast-v0_A2C_lr0.0005_gm0.95

[INFO] Completed run for env=highway-fast-v0, gamma=0.95, lr=0.0005

Final average reward: 8.49


Success rate: 0.30

[INFO] Model saved at: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/models/highway-fast-v0_A2C_lr0.0005_gm0.95_model

2025-03-17 02:36:29.087 python[45071:2415797] +[IMKClient subclass]: chose IMKClient_Modern

2025-03-17 02:36:29.087 python[45071:2415797] +[IMKInputSession subclass]: chose IMKInputSession_Modern

[INFO] Created video: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/videos/highway-fast-v0_A2C_lr0.0005_gm0.95_1742197007.mp4



Index of /

Name	Size	Date Modified
.nofollow/		3/6/25, 4:06:02 AM
.resolve/		3/6/25, 4:06:02 AM
.vol/		3/6/25, 4:06:02 AM
Applications/		3/15/25, 9:12:09 AM
bin/		3/6/25, 4:06:02 AM
cores/		8/4/24, 5:31:41 AM
dev/		3/14/25, 8:24:56 AM
etc/		3/14/25, 8:26:29 AM
home/		3/14/25, 8:26:30 AM
Library/		3/14/25, 8:26:15 AM
opt/		2/10/25, 5:44:30 PM
private/		3/14/25, 8:26:26 AM
sbin/		3/6/25, 4:06:02 AM
System/		3/6/25, 4:06:02 AM
tmp/		3/17/25, 10:21:46 PM
Users/		3/14/25, 8:25:58 AM
usr/		3/6/25, 4:06:02 AM
var/		3/14/25, 8:26:18 AM
Volumes/		3/14/25, 9:18:58 AM
.file	0 B	3/6/25, 4:06:02 AM

TRAINING RUN: highway-fast-v0_A2C_lr0.0005_gm0.8


[INFO] Completed run for env=highway-fast-v0, gamma=0.8, lr=0.0005

Final average reward: 14.10

Success rate: 0.82

[INFO] Model saved at: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/models/highway-fast-v0_A2C_lr0.0005_gm0.8_model

[INFO] Created video: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/videos/highway-fast-v0_A2C_lr0.0005_gm0.8_1742198446.mp4



The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

Reusing TensorBoard on port 6007 (pid 45311), started 0:23:55 ago. (Use '!kill 45311' to kill it.)

Index of /

Name	Size	Date Modified
.nofollow/		3/6/25, 4:06:02 AM
.resolve/		3/6/25, 4:06:02 AM
.vol/		3/6/25, 4:06:02 AM
Applications/		3/15/25, 9:12:09 AM
bin/		3/6/25, 4:06:02 AM
cores/		8/4/24, 5:31:41 AM
dev/		3/14/25, 8:24:56 AM
etc/		3/14/25, 8:26:29 AM
home/		3/14/25, 8:26:30 AM
Library/		3/14/25, 8:26:15 AM
opt/		2/10/25, 5:44:30 PM
private/		3/14/25, 8:26:26 AM
sbin/		3/6/25, 4:06:02 AM
System/		3/6/25, 4:06:02 AM
tmp/		3/17/25, 10:21:46 PM
Users/		3/14/25, 8:25:58 AM
usr/		3/6/25, 4:06:02 AM
var/		3/14/25, 8:26:18 AM
Volumes/		3/14/25, 9:18:58 AM
.file	0 B	3/6/25, 4:06:02 AM

TRAINING RUN: highway-fast-v0_A2C_lr0.001_gm0.95

[INFO] Completed run for env=highway-fast-v0, gamma=0.95, lr=0.001

Final average reward: 14.25

Success rate: 0.25

[INFO] Model saved at: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/models/highway-fast-v0_A2C_lr0.001_gm0.95_model

[INFO] Created video: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/videos/highway-fast-v0_A2C_lr0.001_gm0.95_1742199872.mp4



The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

Reusing TensorBoard on port 6007 (pid 45311), started 0:47:40 ago. (Use '!kill 45311' to kill it.)

Index of /

Name	Size	Date Modified
.nofollow/		3/6/25, 4:06:02 AM
.resolve/		3/6/25, 4:06:02 AM
.vol/		3/6/25, 4:06:02 AM
Applications/		3/15/25, 9:12:09 AM
bin/		3/6/25, 4:06:02 AM
cores/		8/4/24, 5:31:41 AM
dev/		3/14/25, 8:24:56 AM
etc/		3/14/25, 8:26:29 AM
home/		3/14/25, 8:26:30 AM
Library/		3/14/25, 8:26:15 AM
opt/		2/10/25, 5:44:30 PM
private/		3/14/25, 8:26:26 AM
sbin/		3/6/25, 4:06:02 AM
System/		3/6/25, 4:06:02 AM
tmp/		3/17/25, 10:21:46 PM
Users/		3/14/25, 8:25:58 AM
usr/		3/6/25, 4:06:02 AM
var/		3/14/25, 8:26:18 AM
Volumes/		3/14/25, 9:18:58 AM
.file	0 B	3/6/25, 4:06:02 AM

TRAINING RUN: highway-fast-v0_A2C_lr0.001_gm0.8

[INFO] Completed run for env=highway-fast-v0, gamma=0.8, lr=0.001

Final average reward: 7.80

Success rate: 0.48

[INFO] Model saved at: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/models/highway-fast-v0_A2C_lr0.001_gm0.8_model

[INFO] Created video: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/videos/highway-fast-v0_A2C_lr0.001_gm0.8_1742201308.mp4



The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

Reusing TensorBoard on port 6007 (pid 45311), started 1:11:36 ago. (Use '!kill 45311' to kill it.)

Index of /

Name	Size	Date Modified
.nofollow/		3/6/25, 4:06:02 AM
.resolve/		3/6/25, 4:06:02 AM
.vol/		3/6/25, 4:06:02 AM
Applications/		3/15/25, 9:12:09 AM
bin/		3/6/25, 4:06:02 AM
cores/		8/4/24, 5:31:41 AM
dev/		3/14/25, 8:24:56 AM
etc/		3/14/25, 8:26:29 AM
home/		3/14/25, 8:26:30 AM
Library/		3/14/25, 8:26:15 AM
opt/		2/10/25, 5:44:30 PM
private/		3/14/25, 8:26:26 AM
sbin/		3/6/25, 4:06:02 AM
System/		3/6/25, 4:06:02 AM
tmp/		3/17/25, 10:21:46 PM
Users/		3/14/25, 8:25:58 AM
usr/		3/6/25, 4:06:02 AM
var/		3/14/25, 8:26:18 AM
Volumes/		3/14/25, 9:18:58 AM
.file	0 B	3/6/25, 4:06:02 AM

TRAINING RUN: intersection-v0_A2C_lr0.0005_gm0.95

```
/opt/anaconda3/envs/pytorch-env/lib/python3.10/site-packages/gymnasium/envs/registration.py:517: DeprecationWarning: WARN: The environment intersection-v0 is out of date. You should consider upgrading to version `v1`.
```

```
    logger.deprecation(
```

```
[INFO] Completed run for env=intersection-v0, gamma=0.95, lr=0.0005
```

```
    Final average reward: -0.20
```

```
    Success rate: 0.00
```

```
[INFO] Model saved at: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/models/intersection-v0_A2C_lr0.0005_gm0.95_model
```

```
logger.deprecation(
```

A black video player interface. At the bottom, there is a white progress bar. On the left side of the bar is a play button icon, followed by the text '0:00 / 0:20'. To the right of the time are three icons: a speaker icon for audio, a square icon for full screen, and a vertical ellipsis icon for a menu. The rest of the player area is solid black.

```
%reload_ext tensorboard
```

```
Reusing TensorBoard on port 6007 (pid 45311), started 1:53:47 ago. (Use '!kill 45311' to kill it.)
```


Index of /

Name	Size	Date Modified
.nofollow/		3/6/25, 4:06:02 AM
.resolve/		3/6/25, 4:06:02 AM
.vol/		3/6/25, 4:06:02 AM
Applications/		3/15/25, 9:12:09 AM
bin/		3/6/25, 4:06:02 AM
cores/		8/4/24, 5:31:41 AM
dev/		3/14/25, 8:24:56 AM
etc/		3/14/25, 8:26:29 AM
home/		3/14/25, 8:26:30 AM
Library/		3/14/25, 8:26:15 AM
opt/		2/10/25, 5:44:30 PM
private/		3/14/25, 8:26:26 AM
sbin/		3/6/25, 4:06:02 AM
System/		3/6/25, 4:06:02 AM
tmp/		3/17/25, 10:21:46 PM
Users/		3/14/25, 8:25:58 AM
usr/		3/6/25, 4:06:02 AM
var/		3/14/25, 8:26:18 AM
Volumes/		3/14/25, 9:18:58 AM
.file	0 B	3/6/25, 4:06:02 AM

TRAINING RUN: intersection-v0_A2C_lr0.0005_gm0.8

```
/opt/anaconda3/envs/pytorch-env/lib/python3.10/site-packages/gymnasium/envs/registration.py:517: DeprecationWarning: WARN: The environment intersection-v0 is out of date. You should consider upgrading to version `v1`.
```

```
logger.deprecation(
```

```
[INFO] Completed run for env=intersection-v0, gamma=0.8, lr=0.0005
```


```
Final average reward: -0.20
```

```
Success rate: 0.00
```

```
[INFO] Model saved at: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/models/intersection-v0_A2C_lr0.0005_gm0.8_model
```

```
/opt/anaconda3/envs/pytorch-env/lib/python3.10/site-packages/gymnasium/envs/registration.py:517: DeprecationWarning: WARN: The environment intersection-v0 is out of date. You should consider upgrading to version `v1`.
  logger.deprecation(
```

```
[INFO] Created video: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/videos/intersection-v0_A2C_lr0.0005_gm0.8_1742206362.mp4
```



The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

Reusing TensorBoard on port 6007 (pid 45311), started 2:35:51 ago. (Use '!kill 45311' to kill it.)

Index of /

Name	Size	Date Modified
.nofollow/		3/6/25, 4:06:02 AM
.resolve/		3/6/25, 4:06:02 AM
.vol/		3/6/25, 4:06:02 AM
Applications/		3/15/25, 9:12:09 AM
bin/		3/6/25, 4:06:02 AM
cores/		8/4/24, 5:31:41 AM
dev/		3/14/25, 8:24:56 AM
etc/		3/14/25, 8:26:29 AM
home/		3/14/25, 8:26:30 AM
Library/		3/14/25, 8:26:15 AM
opt/		2/10/25, 5:44:30 PM
private/		3/14/25, 8:26:26 AM
sbin/		3/6/25, 4:06:02 AM
System/		3/6/25, 4:06:02 AM
tmp/		3/17/25, 10:21:46 PM
Users/		3/14/25, 8:25:58 AM
usr/		3/6/25, 4:06:02 AM
var/		3/14/25, 8:26:18 AM
Volumes/		3/14/25, 9:18:58 AM
.file	0 B	3/6/25, 4:06:02 AM

TRAINING RUN: intersection-v0_A2C_lr0.001_gm0.95

```
/opt/anaconda3/envs/pytorch-env/lib/python3.10/site-packages/gymnasium/envs/registration.py:517: DeprecationWarning: WARN: The environment intersection-v0 is out of date. You should consider upgrading to version `v1`.
```

```
    logger.deprecation(
```

```
[INF0] Completed run for env=intersection-v0, gamma=0.95, lr=0.001
```


```
    Final average reward: 0.40
```

```
    Success rate: -2.00
```

```
[INF0] Model saved at: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/models/intersection-v0_A2C_lr0.001_gm0.95_model
```

```
/opt/anaconda3/envs/pytorch-env/lib/python3.10/site-packages/gymnasium/envs/registration.py:517: DeprecationWarning: WARN: The environment intersection-v0 is out of date. You should consider upgrading to version `v1`.
  logger.deprecation(
```

```
[INF0] Created video: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/videos/intersection-v0_A2C_lr0.001_gm0.95_1742208888.mp4
```



The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

Reusing TensorBoard on port 6007 (pid 45311), started 3:17:57 ago. (Use '!kill 45311' to kill it.)

Index of /

Name	Size	Date Modified
.nofollow/		3/6/25, 4:06:02 AM
.resolve/		3/6/25, 4:06:02 AM
.vol/		3/6/25, 4:06:02 AM
Applications/		3/15/25, 9:12:09 AM
bin/		3/6/25, 4:06:02 AM
cores/		8/4/24, 5:31:41 AM
dev/		3/14/25, 8:24:56 AM
etc/		3/14/25, 8:26:29 AM
home/		3/14/25, 8:26:30 AM
Library/		3/14/25, 8:26:15 AM
opt/		2/10/25, 5:44:30 PM
private/		3/14/25, 8:26:26 AM
sbin/		3/6/25, 4:06:02 AM
System/		3/6/25, 4:06:02 AM
tmp/		3/17/25, 10:21:46 PM
Users/		3/14/25, 8:25:58 AM
usr/		3/6/25, 4:06:02 AM
var/		3/14/25, 8:26:18 AM
Volumes/		3/14/25, 9:18:58 AM
.file	0 B	3/6/25, 4:06:02 AM

TRAINING RUN: intersection-v0_A2C_lr0.001_gm0.8

```
/opt/anaconda3/envs/pytorch-env/lib/python3.10/site-packages/gymnasium/envs/registration.py:517: DeprecationWarning: WARN: The environment intersection-v0 is out of date. You should consider upgrading to version `v1`.
```

```
    logger.deprecation(
```

```
[INFO] Completed run for env=intersection-v0, gamma=0.8, lr=0.001
```


```
    Final average reward: -0.20
```

```
    Success rate: 0.00
```

```
[INFO] Model saved at: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/models/intersection-v0_A2C_lr0.001_gm0.8_model
```

```
/opt/anaconda3/envs/pytorch-env/lib/python3.10/site-packages/gymnasium/envs/registration.py:517: DeprecationWarning: WARN: The environment intersection-v0 is out of date. You should consider upgrading to version `v1`.
  logger.deprecation(
```

```
[INFO] Created video: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/videos/intersection-v0_A2C_lr0.001_gm0.8_1742211425.mp4
```



The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

Reusing TensorBoard on port 6007 (pid 45311), started 4:00:14 ago. (Use '!kill 45311' to kill it.)

Index of /

Name	Size	Date Modified
.nofollow/		3/6/25, 4:06:02 AM
.resolve/		3/6/25, 4:06:02 AM
.vol/		3/6/25, 4:06:02 AM
Applications/		3/15/25, 9:12:09 AM
bin/		3/6/25, 4:06:02 AM
cores/		8/4/24, 5:31:41 AM
dev/		3/14/25, 8:24:56 AM
etc/		3/14/25, 8:26:29 AM
home/		3/14/25, 8:26:30 AM
Library/		3/14/25, 8:26:15 AM
opt/		2/10/25, 5:44:30 PM
private/		3/14/25, 8:26:26 AM
sbin/		3/6/25, 4:06:02 AM
System/		3/6/25, 4:06:02 AM
tmp/		3/17/25, 10:21:46 PM
Users/		3/14/25, 8:25:58 AM
usr/		3/6/25, 4:06:02 AM
var/		3/14/25, 8:26:18 AM
Volumes/		3/14/25, 9:18:58 AM
.file	0 B	3/6/25, 4:06:02 AM

TRAINING RUN: roundabout-v0_A2C_lr0.0005_gm0.95


[INFO] Completed run for env=roundabout-v0, gamma=0.95, lr=0.0005

Final average reward: 6.73

Success rate: 0.72

[INFO] Model saved at: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/models/roundabout-v0_A2C_lr0.0005_gm0.95_model

[INFO] Created video: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/videos/roundabout-v0_A2C_lr0.0005_gm0.95_1742212954.mp4



The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

Reusing TensorBoard on port 6007 (pid 45311), started 4:25:43 ago. (Use '!kill 45311' to kill it.)

Index of /

Name	Size	Date Modified
.nofollow/		3/6/25, 4:06:02 AM
.resolve/		3/6/25, 4:06:02 AM
.vol/		3/6/25, 4:06:02 AM
Applications/		3/15/25, 9:12:09 AM
bin/		3/6/25, 4:06:02 AM
cores/		8/4/24, 5:31:41 AM
dev/		3/14/25, 8:24:56 AM
etc/		3/14/25, 8:26:29 AM
home/		3/14/25, 8:26:30 AM
Library/		3/14/25, 8:26:15 AM
opt/		2/10/25, 5:44:30 PM
private/		3/14/25, 8:26:26 AM
sbin/		3/6/25, 4:06:02 AM
System/		3/6/25, 4:06:02 AM
tmp/		3/17/25, 10:21:46 PM
Users/		3/14/25, 8:25:58 AM
usr/		3/6/25, 4:06:02 AM
var/		3/14/25, 8:26:18 AM
Volumes/		3/14/25, 9:18:58 AM
.file	0 B	3/6/25, 4:06:02 AM

TRAINING RUN: roundabout-v0_A2C_lr0.0005_gm0.8


[INFO] Completed run for env=roundabout-v0, gamma=0.8, lr=0.0005

Final average reward: 4.25

Success rate: 0.33

[INFO] Model saved at: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/models/roundabout-v0_A2C_lr0.0005_gm0.8_model

[INFO] Created video: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/videos/roundabout-v0_A2C_lr0.0005_gm0.8_1742214479.mp4



The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

Reusing TensorBoard on port 6007 (pid 45311), started 4:51:08 ago. (Use '!kill 45311' to kill it.)

Index of /

Name	Size	Date Modified
.nofollow/		3/6/25, 4:06:02 AM
.resolve/		3/6/25, 4:06:02 AM
.vol/		3/6/25, 4:06:02 AM
Applications/		3/15/25, 9:12:09 AM
bin/		3/6/25, 4:06:02 AM
cores/		8/4/24, 5:31:41 AM
dev/		3/14/25, 8:24:56 AM
etc/		3/14/25, 8:26:29 AM
home/		3/14/25, 8:26:30 AM
Library/		3/14/25, 8:26:15 AM
opt/		2/10/25, 5:44:30 PM
private/		3/14/25, 8:26:26 AM
sbin/		3/6/25, 4:06:02 AM
System/		3/6/25, 4:06:02 AM
tmp/		3/17/25, 10:21:46 PM
Users/		3/14/25, 8:25:58 AM
usr/		3/6/25, 4:06:02 AM
var/		3/14/25, 8:26:18 AM
Volumes/		3/14/25, 9:18:58 AM
.file	0 B	3/6/25, 4:06:02 AM

TRAINING RUN: roundabout-v0_A2C_lr0.001_gm0.95


[INFO] Completed run for env=roundabout-v0, gamma=0.95, lr=0.001

Final average reward: 6.33

Success rate: 0.58

[INFO] Model saved at: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/models/roundabout-v0_A2C_lr0.001_gm0.95_model

[INFO] Created video: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/videos/roundabout-v0_A2C_lr0.001_gm0.95_1742216015.mp4



The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

Reusing TensorBoard on port 6007 (pid 45311), started 5:16:44 ago. (Use '!kill 45311' to kill it.)

Index of /

Name	Size	Date Modified
.nofollow/		3/6/25, 4:06:02 AM
.resolve/		3/6/25, 4:06:02 AM
.vol/		3/6/25, 4:06:02 AM
Applications/		3/15/25, 9:12:09 AM
bin/		3/6/25, 4:06:02 AM
cores/		8/4/24, 5:31:41 AM
dev/		3/14/25, 8:24:56 AM
etc/		3/14/25, 8:26:29 AM
home/		3/14/25, 8:26:30 AM
Library/		3/14/25, 8:26:15 AM
opt/		2/10/25, 5:44:30 PM
private/		3/14/25, 8:26:26 AM
sbin/		3/6/25, 4:06:02 AM
System/		3/6/25, 4:06:02 AM
tmp/		3/17/25, 10:21:46 PM
Users/		3/14/25, 8:25:58 AM
usr/		3/6/25, 4:06:02 AM
var/		3/14/25, 8:26:18 AM
Volumes/		3/14/25, 9:18:58 AM
.file	0 B	3/6/25, 4:06:02 AM

TRAINING RUN: roundabout-v0_A2C_lr0.001_gm0.8


[INFO] Completed run for env=roundabout-v0, gamma=0.8, lr=0.001

Final average reward: 5.20

Success rate: 0.42

[INFO] Model saved at: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/models/roundabout-v0_A2C_lr0.001_gm0.8_model

[INFO] Created video: /Users/sanjayuppala/Desktop/CSC 580 AI_2/Project_A2C/Project_A2C/videos/roundabout-v0_A2C_lr0.001_gm0.8_1742217545.mp4



The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

Reusing TensorBoard on port 6007 (pid 45311), started 5:42:14 ago. (Use '!kill 45311' to kill it.)

Index of /

Name	Size	Date Modified
.nofollow/		3/6/25, 4:06:02 AM
.resolve/		3/6/25, 4:06:02 AM
.vol/		3/6/25, 4:06:02 AM
Applications/		3/15/25, 9:12:09 AM
bin/		3/6/25, 4:06:02 AM
cores/		8/4/24, 5:31:41 AM
dev/		3/14/25, 8:24:56 AM
etc/		3/14/25, 8:26:29 AM
home/		3/14/25, 8:26:30 AM
Library/		3/14/25, 8:26:15 AM
opt/		2/10/25, 5:44:30 PM
private/		3/14/25, 8:26:26 AM
sbin/		3/6/25, 4:06:02 AM
System/		3/6/25, 4:06:02 AM
tmp/		3/17/25, 10:21:46 PM
Users/		3/14/25, 8:25:58 AM
usr/		3/6/25, 4:06:02 AM
var/		3/14/25, 8:26:18 AM
Volumes/		3/14/25, 9:18:58 AM
.file	0 B	3/6/25, 4:06:02 AM

Final Result of All model trained Above


highway-fast-v0_A2C_lr0.0005_gm0.95
highway-fast-v0_A2C_lr0.0005_gm0.8
highway-fast-v0_A2C_lr0.001_gm0.95
highway-fast-v0_A2C_lr0.001_gm0.8
intersection-v0_A2C_lr0.0005_gm0.95
intersection-v0_A2C_lr0.0005_gm0.8
intersection-v0_A2C_lr0.001_gm0.95
intersection-v0_A2C_lr0.001_gm0.8
roundabout-v0_A2C_lr0.0005_gm0.95
roundabout-v0_A2C_lr0.0005_gm0.8
roundabout-v0_A2C_lr0.001_gm0.95
roundabout-v0_A2C_lr0.001_gm0.8

```
In [10]: #To kill the prvious tensboard
os.system("pkill -f tensorboard")
# Display TENSORBOARD plot for this run
if IN_JUPYTER:
    get_ipython().run_line_magic('load_ext', 'tensorboard')
    get_ipython().run_line_magic('tensorboard', '--logdir ./Project_A2C/logs')
```

The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard

Index of /

Name	Size	Date Modified
.nofollow/		3/6/25, 4:06:02 AM
.resolve/		3/6/25, 4:06:02 AM
.vol/		3/6/25, 4:06:02 AM
Applications/		3/15/25, 9:12:09 AM
bin/		3/6/25, 4:06:02 AM
cores/		8/4/24, 5:31:41 AM
dev/		3/14/25, 8:24:56 AM
etc/		3/14/25, 8:26:29 AM
home/		3/14/25, 8:26:30 AM
Library/		3/14/25, 8:26:15 AM
opt/		2/10/25, 5:44:30 PM
private/		3/14/25, 8:26:26 AM
sbin/		3/6/25, 4:06:02 AM
System/		3/6/25, 4:06:02 AM
tmp/		3/17/25, 10:21:46 PM
Users/		3/14/25, 8:25:58 AM
usr/		3/6/25, 4:06:02 AM
var/		3/14/25, 8:26:18 AM
Volumes/		3/14/25, 9:18:58 AM
.file	0 B	3/6/25, 4:06:02 AM

 TensorBoard is running at http://localhost:6015/

```
In [ ]:
```